

Uploading Kaggle Json file

In []:

```
#we will upload kaggle.json file which will help us in loading kaggle dataset here in colab.
from google.colab import files
files.upload()

# making kaggle directory and copying kaggle.json file inside kaggle directory.
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

Choose File No file selected

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

Downloading Pretrained Weights

In []:

```
!kaggle datasets download -d absolutegaming/pre-trained-weights #Load Pretrained Weights
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /root/.kaggle/kaggle.json'

Downloading pre-trained-weights.zip to /content

98% 216M/220M [00:03<00:00, 49.3MB/s]

100% 220M/220M [00:03<00:00, 66.4MB/s]

In []:

```
!unzip '/content/pre-trained-weights.zip'
```

Archive: /content/pre-trained-weights.zip
inflating: yolov3.weights

Downloading main data

In []:

```
!wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.4147.135 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-US,en;q=0.9" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/19989/1160143/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1598706338&Signature=eKQ5C%2FDTJiP6YJ10aKXxMMcXpw7SpppFE1w1%2B9DEdsQ10yCsJlbADBuVfBAVDylah6JqmcKas8i%2FKPNC9spRd8QvFySBlwd6UcWZUXz%2BJQLf9BM3Ez1OHjweVOnV2hYb7H0LUqUC44Upd2aAtJhvJ7NDamCheva7iTyMh%2F24au78nytB2oVfshyag3QYNg8bxW7nT7Dy3A6DU9jFFGr1TuKHD%2FS0wHscglvehA%2FEicH0Dr%2Bd7yXEZaAfGm%2BGzk%2FC5%2FxFRNdFco%2BzSCFLQcYpZg8E5ceaFgDDvtNbWGvB1yP3MTneCAIRZNrAjheZtWKaPK43Q%3D%3D&response-content-disposition=attachment%3B+filename%3Dglobal-wheat-detection.zip" -c -O 'global-wheat-detection.zip'
```

```
--2020-08-27 05:37:36-- https://storage.googleapis.com/kaggle-competitions-data/kaggle-v2/19989/1160143/bundle/archive.zip?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1598706338&Signature=eKQ5C%2FDTJiP6YJ10aKXxMMcXpw7SpppFE1w1%2B9DEdsQ10yCsJlbADBuVfBAVDylah6JqmcKas8i%2FKPNC9spRd8QvFySBlwd6UcWZUXz%2BJQLf9BM3Ez1OHjweVOnV2hYb7H0LUqUC44Upd2aAtJhvJ7NDamCheva7iTyMh%2F24au78nytB2oVfshyag3QYNg8bxW7nT7Dy3A6DU9jFFGr1TuKHD%2FS0wHscglvehA%2FEicH0Dr%2Bd7yXEZaAfGm%2BGzk%2FC5%2FxFRNdFco%2BzSCFLQcYpZg8E5ceaFgDDvtNbWGvB1yP3MTneCAIRZNrAjheZtWKaPK43Q%3D%3D&response-content-disposition=attachment%3B+filename%3Dglobal-wheat-detection.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.142.128, 74.125.145.128
```

```
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.142.128, 74.125.193.128, 74.125.20.128, ...  
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.142.128|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 636694264 (607M) [application/zip]  
Saving to: 'global-wheat-detection.zip'
```

```
global-wheat-detect 100%[=====>] 607.20M 76.8MB/s in 8.6s
```

```
2020-08-27 05:37:45 (70.6 MB/s) - 'global-wheat-detection.zip' saved [636694264/636694264]
```



```
In [ ]:
```

```
!unzip 'global-wheat-detection.zip'
```

Loading all the necessary Modules

```
In [ ]:
```

```
from absl import logging  
import tensorflow as tf  
from tensorflow.keras import Model  
from tensorflow.keras.layers import  
Add, Concatenate, Conv2D, Input, Lambda, LeakyReLU, MaxPool2D, UpSampling2D, ZeroPadding2D, BatchNormalizati  
on  
from tensorflow.keras.regularizers import l2  
from tensorflow.keras.losses import binary_crossentropy, sparse_categorical_crossentropy  
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint, TensorBoard  
from tensorflow.keras.callbacks import ReduceLROnPlateau, ModelCheckpoint, EarlyStopping,  
LearningRateScheduler, CSVLogger  
import seaborn as sns  
import matplotlib.pyplot as plt  
import tensorflow as tf  
import numpy as np  
from PIL import Image, ImageDraw, ImageFont  
from IPython.display import display  
from seaborn import color_palette  
import cv2  
import pandas as pd  
from tqdm import tqdm  
import os  
import tensorflow as tf  
import time  
from tqdm import tqdm  
import datetime  
import ast
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning:  
pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.  
import pandas.util.testing as tm
```

Loading Data into Dataframe

```
In [ ]:
```

```
train_df = pd.read_csv('/content/train.csv') #Load the CSV file  
train_df['image_id'] = '/content/train/' + train_df['image_id'].astype(str) + '.jpg' #Add the path to  
the images
```

Preprocessing the data

```
In [ ]:
```

```
train_df[['x_min', 'y_min', 'width', 'height']] = pd.DataFrame([ast.literal_eval(x) for x in train_d  
f.bbox.tolist()], index= train_df.index)  
train_df = train_df[['image_id', 'bbox', 'source', 'x_min', 'y_min', 'width', 'height']]  
train_df['area'] = train_df['width'] * train_df['height']
```

```

train_df['x_max'] = train_df['x_min'] + train_df['width']
train_df['y_max'] = train_df['y_min'] + train_df['height']
train_df = train_df.drop(['bbox', 'source'], axis=1)
train_df = train_df[['image_id', 'x_min', 'y_min', 'x_max', 'y_max', 'width', 'height', 'area']]

# There are some buggy annotations in training images having huge bounding boxes. Let's remove those
train_df = train_df[train_df['area'] < 100000]

```

In []:

```

train_df = train_df.drop(['width', 'height', 'area'], axis=1) #Dropping unwanted columns
class_list = ['Wheat']*train_df.shape[0]
train_df['category'] = class_list #Adding category column
label_list = [1]*train_df.shape[0]
train_df['label'] = label_list #Adding class label to each row
train_df.columns = ['img_path', 'min_w', 'min_h', 'max_w', 'max_h', 'category', 'label'] #Changing Column names

df = train_df #Loading a copy into df
df = df[['max_h', 'max_w', 'min_h', 'min_w', 'category', 'img_path', 'label']] #Changing the order of columns
print(df.head(5))

```

	max_h	max_w	min_h	min_w	category	img_path	label
0	258.0	890.0	222.0	834.0	Wheat	/content/train/b6ab77fd7.jpg	1
1	606.0	356.0	548.0	226.0	Wheat	/content/train/b6ab77fd7.jpg	1
2	664.0	451.0	504.0	377.0	Wheat	/content/train/b6ab77fd7.jpg	1
3	202.0	943.0	95.0	834.0	Wheat	/content/train/b6ab77fd7.jpg	1
4	261.0	150.0	144.0	26.0	Wheat	/content/train/b6ab77fd7.jpg	1

Test Train Split

In []:

```

max_boxes = df.groupby(['img_path']).count()['category'].max()
print('Max number of boxes in a given image are :', max_boxes) #Calculating maximum number of boxes in an image

index = list(set([i.split('/')[ -1] for i in df['img_path'].values]))
print('Total number of unique images in the dataframe :', len(index)) #Taking out all the unique images in the df

#Here for simplicity sake i am using 1000 images as train and 200 images as test.
print('Splitting into train and test data.....')
train_image=index[0:2500]
test_image =index[2500:3733]
print('Number of images in train data :', len(train_image))
print('Number of images in test data :', len(test_image))

```

```

Max number of boxes in a given image are : 116
Total number of unique images in the dataframe : 3733
Splitting into train and test data.....
Number of images in train data : 2500
Number of images in test data : 873

```

Initializing all the necessary variables

In []:

```

yolo_anchors = np.array([(10, 13), (16, 30), (33, 23), (30, 61), (62, 45),
                          (59, 119), (116, 90), (156, 198), (373, 326)],
                          np.float32) / 320
yolo_anchor_masks = np.array([[6, 7, 8], [3, 4, 5], [0, 1, 2]])

size = 320 # size of resize image
batch_size = 8 #Batch size that we want to load when training
yolo_max_boxes = 116 # maximum yolo boxes predicted per image.
yolo_iou_threshold = 0.5 # IOU threshold score
yolo_score_threshold = 0.4 # objectness threshold score

```

```

learning_rate = 1e-4 # learning_rate
num_classes = 2 # num of category in our dataset
epochs = 30 # epochs run to fine tune our model
class_dict = {'Wheat':1} # Mapping of category and its corresponding label

```

YOLOV3 architecture starts here

In []:

```

def DarknetConv(x, filters, size, strides=1, batch_norm=True):
    if strides == 1:
        padding = 'same'
    else:
        x = ZeroPadding2D((1, 0), (1, 0))(x) # top left half-padding
        padding = 'valid'
    x = Conv2D(filters=filters, kernel_size=size,
               strides=strides, padding=padding,
               use_bias=not batch_norm, kernel_regularizer=l2(0.0005))(x)
    if batch_norm:
        x = BatchNormalization()(x)
        x = LeakyReLU(alpha=0.1)(x) # by default alpha is 0.3
    return x

```

In []:

```

# below is residual connection layer function
def DarknetResidual(x, filters):
    prev = x # storing input in prev variable
    x = DarknetConv(x, filters // 2, 1)
    x = DarknetConv(x, filters, 3)
    x = Add()([prev, x]) # residual connection
    return x

```

In []:

```

# this is our real Darknetblock function calling above 2 fucntions
def DarknetBlock(x, filters, blocks):
    x = DarknetConv(x, filters, 3, strides=2)
    for _ in range(blocks):
        x = DarknetResidual(x, filters)
    return x

```

In []:

```

def Darknet(name=None):
    x = inputs = Input([None, None, 3])
    x = DarknetConv(x, 32, 3)
    x = DarknetBlock(x, 64, 1)
    x = DarknetBlock(x, 128, 2) # skip connection
    x = x_36 = DarknetBlock(x, 256, 8) # skip connection
    x = x_61 = DarknetBlock(x, 512, 8)
    x = DarknetBlock(x, 1024, 4) # last layer detecting bounding box dimension (tx,ty,bx,by)
    return tf.keras.Model(inputs, (x_36, x_61, x), name=name)

```

In []:

```

def YoloConv(filters, name=None):
    def yolo_conv(x_in):
        if isinstance(x_in, tuple):
            inputs = Input(x_in[0].shape[1:]), Input(x_in[1].shape[1:])
            x, x_skip = inputs

            # concat with skip connection
            x = DarknetConv(x, filters, 1)
            # upsampling of a layer
            x = UpSampling2D(2)(x)
            # concatenation of skip connection result and last output result
            x = Concatenate()([x, x_skip])
        else:
            x = inputs = Input(x_in.shape[1:])

```

```

x = DarknetConv(x, filters, 1)
x = DarknetConv(x, filters * 2, 3)
x = DarknetConv(x, filters, 1)
x = DarknetConv(x, filters * 2, 3)
x = DarknetConv(x, filters, 1)
    return Model(inputs, x, name=name)(x_in)
return yolo_conv

```

YOLO output

In []:

```

def YoloOutput(filters, anchors, classes, name=None):
    def yolo_output(x_in):
        x = inputs = Input(x_in.shape[1:]) # this is an input shape excluded with batch size.
        x = DarknetConv(x, filters * 2, 3) # Darkconv is a fn implemented above which is internally
calling
        x = DarknetConv(x, anchors * (classes + 5), 1, batch_norm=False)
        x = Lambda(lambda x: tf.reshape(x, (-1, tf.shape(x)[1], tf.shape(x)[2],
            anchors, classes + 5)))(x)
        # x is reshaped into (None, grid_size, grid_size, anchors, (x,y,w,h,objectness
score,..classes))
        return tf.keras.Model(inputs, x, name=name)(x_in)
    return yolo_output

```

Non Max Supression

In []:

```

def yolo_boxes(pred, anchors, classes):
    # pred: (batch_size, grid, grid, anchors, (x, y, w, h, obj, ...classes))
    grid_size = tf.shape(pred)[1]
    box_xy, box_wh, objectness, class_probs = tf.split(
        pred, (2, 2, 1, classes), axis=-1)
    # objectness : it's means whether there is any object in a predicted box
    # class_probs : it's a probability of a class given object is there i.e P(Pc|object)

    box_xy = tf.sigmoid(box_xy)
    objectness = tf.sigmoid(objectness)
    class_probs = tf.sigmoid(class_probs)
    pred_box = tf.concat([box_xy, box_wh], axis=-1) # original xywh for loss

    # !!! grid[x][y] == (y, x)
    grid = tf.meshgrid(tf.range(grid_size), tf.range(grid_size))
    grid = tf.expand_dims(tf.stack(grid, axis=-1), axis=2) # [gx, gy, 1, 2]
    box_xy = (box_xy + tf.cast(grid, tf.float32)) / \
        tf.cast(grid_size, tf.float32)
    box_wh = tf.exp(box_wh) * anchors

    box_x1y1 = box_xy - box_wh / 2
    box_x2y2 = box_xy + box_wh / 2
    bbox = tf.concat([box_x1y1, box_x2y2], axis=-1)

    return bbox, objectness, class_probs, pred_box

def yolo_nms(outputs, anchors, masks, classes):
    # boxes, confidence scores(objectness scores), class probabilities

    b, c, t = [], [], []

    # iterating through each outputs predicted by model
    for o in outputs:
        b.append(tf.reshape(o[0], (tf.shape(o[0])[0], -1, tf.shape(o[0])[-1])))
        c.append(tf.reshape(o[1], (tf.shape(o[1])[0], -1, tf.shape(o[1])[-1])))
        t.append(tf.reshape(o[2], (tf.shape(o[2])[0], -1, tf.shape(o[2])[-1])))

    bbox = tf.concat(b, axis=1)
    confidence = tf.concat(c, axis=1)
    class_probs = tf.concat(t, axis=1)

    scores = confidence * class_probs # this is P(Pc|objectness score) value

```

```

boxes, scores, classes, valid_detections = tf.image.combined_non_max_suppression(
    boxes=tf.reshape(bbox, (tf.shape(bbox)[0], -1, 1, 4)),
    scores=tf.reshape(
        scores, (tf.shape(scores)[0], -1, tf.shape(scores)[-1])),
    max_output_size_per_class = yolo_max_boxes, # here it is 223, define above
    max_total_size = yolo_max_boxes,
    iou_threshold = yolo_iou_threshold, # threshold for filtering the boxes
    score_threshold = yolo_score_threshold # threshold for objectness score below which we ignore that bounding box
)

return boxes, scores, classes, valid_detections

```

Backbone Darknet Architecture

In []:

```

def YoloV3(size=None, channels=3, anchors=yolo_anchors,
           masks=yolo_anchor_masks, classes=80, training=False):
    x = inputs = Input([size, size, channels], name='input') # input of an image

    x_36, x_61, x = Darknet(name='yolo_darknet')(x) # backbone networks 3 outputs w.r.t to each grid size
    # till here darknet network

    # from below it's a Feature Pyramid Network with lateral connections
    # for 13*13 grid size output
    x = YoloConv(512, name='yolo_conv_0')(x)
    output_0 = YoloOutput(512, len(masks[0]), classes, name='yolo_output_0')(x)

    # for 26*26 grid size output
    x = YoloConv(256, name='yolo_conv_1')((x, x_61))
    output_1 = YoloOutput(256, len(masks[1]), classes, name='yolo_output_1')(x)

    # for 52*52 grid size output
    x = YoloConv(128, name='yolo_conv_2')((x, x_36))
    output_2 = YoloOutput(128, len(masks[2]), classes, name='yolo_output_2')(x)

    if training:
        return Model(inputs, (output_0, output_1, output_2), name='yolov3')

    # for 13*13 grid size output
    boxes_0 = Lambda(lambda x: yolo_boxes(x, anchors[masks[0]], classes),
                     name='yolo_boxes_0')(output_0)
    # for 26*26 grid size output
    boxes_1 = Lambda(lambda x: yolo_boxes(x, anchors[masks[1]], classes),
                     name='yolo_boxes_1')(output_1)
    # for 52*52 grid size output
    boxes_2 = Lambda(lambda x: yolo_boxes(x, anchors[masks[2]], classes),
                     name='yolo_boxes_2')(output_2)
    # after combining boxes from various scales we have total 10,647 boxes which is too large
    # so to remove invalid boxes we use non_maximum_suppression

    outputs = Lambda(lambda x: yolo_nms(x, anchors, masks, classes),
                     name='yolo_nms')((boxes_0[:3], boxes_1[:3], boxes_2[:3]))

    return Model(inputs, outputs, name='yolov3')

```

Custom Loss Function

In []:

```

def YoloLoss(anchors, classes=80, ignore_thresh=0.5):
    def yolo_loss(y_true, y_pred):
        # 1. transform all pred outputs
        # y_pred: (batch_size, grid, grid, anchors, (x, y, w, h, obj, ...cls))
        pred_box, pred_obj, pred_class, pred_xywh = yolo_boxes(
            y_pred, anchors, classes)
        pred_xy = pred_xywh[..., 0:2]
        pred_wh = pred_xywh[..., 2:4]

        # 2. transform all true outputs

```

```

# 2. transform all true outputs
# y_true: (batch_size, grid, grid, anchors, (x1, y1, x2, y2, obj, cls))

true_box, true_obj, true_class_idx = tf.split(
    y_true, (4, 1, 1), axis=-1)

# the above split function split (x1,y1,x2...cls) into (x1,y1),(x2,y2),(obj),(cls)
# the 4,1,1 is a length at which it split

true_xy = (true_box[..., 0:2] + true_box[..., 2:4]) / 2 # finding center (Xcen,Ycen)
true_wh = true_box[..., 2:4] - true_box[..., 0:2] # width and height

# give higher weights to small boxes
box_loss_scale = 2 - true_wh[..., 0] * true_wh[..., 1]

# 3. inverting the pred box equations
grid_size = tf.shape(y_true)[1]
grid = tf.meshgrid(tf.range(grid_size), tf.range(grid_size))
grid = tf.expand_dims(tf.stack(grid, axis=-1), axis=2)
true_xy = true_xy * tf.cast(grid_size, tf.float32) - \
    tf.cast(grid, tf.float32) # this code snippet giving us at which point each cell is starting and ending

# in resize image of 416 * 416
# suppose there 13*13 = 169 cells , so every cell we will have starting and ending point

true_wh = tf.math.log(true_wh / anchors)
# YOLO doesn't predict the absolute coordinates of the bounding box's center
true_wh = tf.where(tf.math.is_inf(true_wh),
    tf.zeros_like(true_wh), true_wh)

# 4. calculate all masks
obj_mask = tf.squeeze(true_obj, -1)

# ignore false positive when iou is over threshold
best_iou = tf.map_fn(
    lambda x: tf.reduce_max(broadcast_iou(x[0], tf.boolean_mask(
        x[1], tf.cast(x[2], tf.bool))), axis=-1),
    (pred_box, true_box, obj_mask),
    tf.float32)
ignore_mask = tf.cast(best_iou < ignore_thresh, tf.float32)

# 5. calculate all losses
xy_loss = obj_mask * box_loss_scale * \
    tf.reduce_sum(tf.square(true_xy - pred_xy), axis=-1)
wh_loss = obj_mask * box_loss_scale * \
    tf.reduce_sum(tf.square(true_wh - pred_wh), axis=-1)

obj_loss = binary_crossentropy(true_obj, pred_obj)
# obj_loss = obj_mask * obj_loss + \
#     (1 - obj_mask) * ignore_mask * obj_loss

alpha = 0.75 # focal loss hyperparameter
conf_focal = tf.pow(obj_mask-tf.squeeze(tf.sigmoid(pred_obj),-1),2)
obj_loss = conf_focal*((1-alpha)*obj_mask*obj_loss +
alpha*(1-obj_mask)*ignore_mask*obj_loss) # batch * grid * grid * anchors_per_scale

# TODO: use binary_crossentropy instead
class_loss = obj_mask * sparse_categorical_crossentropy(
    true_class_idx, pred_class)

# 6. sum over (batch, gridx, gridy, anchors) => (batch, 1)
xy_loss = tf.reduce_sum(xy_loss, axis=(1, 2, 3))
wh_loss = tf.reduce_sum(wh_loss, axis=(1, 2, 3))
obj_loss = tf.reduce_sum(obj_loss, axis=(1, 2, 3))
class_loss = tf.reduce_sum(class_loss, axis=(1, 2, 3))

return xy_loss + wh_loss + obj_loss + class_loss
return yolo_loss

```

Some useful functions

In []:

```

YOLOV3_LAYER_LIST = ['yolo_darknet','yolo_conv_0','yolo_output_0','yolo_conv_1',
    'yolo output 1','yolo conv 2','yolo output 2',]

```

```

def load_darknet_weights(model, weights_file, tiny=False):
    wf = open(weights_file, 'rb') # reading weights file
    major, minor, revision, seen, _ = np.fromfile(wf, dtype=np.int32, count=5)

    layers = YOLOV3_LAYER_LIST

    # iterating through all layers define in above yolov3_layers_list
    for layer_name in layers:
        # for eg if there is one layer darknet then there is many sub layers inside it's network
        sub_model = model.get_layer(layer_name)
        for i, layer in enumerate(sub_model.layers):
            if not layer.name.startswith('conv2d'):
                continue
            batch_norm = None
            if i + 1 < len(sub_model.layers) and \
                sub_model.layers[i + 1].name.startswith('batch_norm'):
                batch_norm = sub_model.layers[i + 1]

            logging.info("{} / {} {}".format(
                sub_model.name, layer.name, 'bn' if batch_norm else 'bias'))

            filters = layer.filters
            size = layer.kernel_size[0]
            in_dim = layer.input_shape[-1]

            if batch_norm is None:
                conv_bias = np.fromfile(wf, dtype=np.float32, count=filters)
            else:
                # darknet [beta, gamma, mean, variance]
                bn_weights = np.fromfile(
                    wf, dtype=np.float32, count=4 * filters)
                # tf [gamma, beta, mean, variance]
                bn_weights = bn_weights.reshape((4, filters))[[1, 0, 2, 3]]

            # darknet shape (out_dim, in_dim, height, width)
            conv_shape = (filters, in_dim, size, size)
            conv_weights = np.fromfile(
                wf, dtype=np.float32, count=np.product(conv_shape))
            # tf shape (height, width, in_dim, out_dim)
            conv_weights = conv_weights.reshape(
                conv_shape).transpose([2, 3, 1, 0])

            if batch_norm is None:
                layer.set_weights([conv_weights, conv_bias])
            else:
                layer.set_weights([conv_weights])
                batch_norm.set_weights(bn_weights)

    assert len(wf.read()) == 0, 'failed to read all data'
    wf.close()

def broadcast_iou(box_1, box_2):
    # box_1: (... , (x1, y1, x2, y2))
    # box_2: (N, (x1, y1, x2, y2))

    # broadcast boxes
    box_1 = tf.expand_dims(box_1, -2)
    box_2 = tf.expand_dims(box_2, 0)
    # new_shape: (... , N, (x1, y1, x2, y2))
    new_shape = tf.broadcast_dynamic_shape(tf.shape(box_1), tf.shape(box_2))
    box_1 = tf.broadcast_to(box_1, new_shape) # it will change the shape of box into new_shape
    given
    box_2 = tf.broadcast_to(box_2, new_shape)

    # in below code we are finding intersection box width and height through which we will find in
    # intersection area.
    # and this we are finding all boxes
    int_w = tf.maximum(tf.minimum(box_1[..., 2], box_2[..., 2]) -
                        tf.maximum(box_1[..., 0], box_2[..., 0]), 0)
    int_h = tf.maximum(tf.minimum(box_1[..., 3], box_2[..., 3]) -
                        tf.maximum(box_1[..., 1], box_2[..., 1]), 0)
    int_area = int_w * int_h # area of intersection
    box_1_area = (box_1[..., 2] - box_1[..., 0]) * \
        (box_1[..., 3] - box_1[..., 1]) # this box 1 area contains all boxes area predicted in an i

```



```

image
box_2_area = (box_2[..., 2] - box_2[..., 0]) * \
              (box_2[..., 3] - box_2[..., 1]) # this box2_area is our ground truth box area

# Formula: Union(A,B) = A + B - Inter(A,B)
return int_area / (box_1_area + box_2_area - int_area)

def freeze_all(model, frozen=True):
    model.trainable = not frozen
    if isinstance(model, tf.keras.Model):
        for l in model.layers:
            freeze_all(l, frozen)

```

Dataset Loader And Transformation

In []:

```

@tf.function
def transform_targets_for_output(y_true, grid_size, anchor_idxs):
    # y_true: (N, boxes, (x1, y1, x2, y2, class, best_anchor))
    N = tf.shape(y_true)[0]

    # y_true_out: (N, grid, grid, anchors, [x, y, w, h, obj, class])
    y_true_out = tf.zeros(
        (N, grid_size, grid_size, tf.shape(anchor_idxs)[0], 6))

    anchor_idxs = tf.cast(anchor_idxs, tf.int32)

    indexes = tf.TensorArray(tf.int32, 1, dynamic_size=True)
    updates = tf.TensorArray(tf.float32, 1, dynamic_size=True)
    idx = 0

    # below iteration change the values and update it to the format which acceptable by yolov3.
    for i in tf.range(N):
        for j in tf.range(tf.shape(y_true)[1]):
            if tf.equal(y_true[i][j][2], 0):
                continue
            anchor_eq = tf.equal(
                anchor_idxs, tf.cast(y_true[i][j][5], tf.int32))

            if tf.reduce_any(anchor_eq):
                box = y_true[i][j][0:4] # (x1,y1,x2,y2)
                box_xy = (y_true[i][j][0:2] + y_true[i][j][2:4]) / 2 # ((x1+x2)/2, (y1+y2)/2)

                # which is (Xcenter,Ycenter)

                anchor_idx = tf.cast(tf.where(anchor_eq), tf.int32)
                grid_xy = tf.cast(box_xy // (1/grid_size), tf.int32) # multiplying it by grid_size

                # grid[y][x][anchor] = (tx, ty, bw, bh, obj, class)
                indexes = indexes.write(
                    idx, [i, grid_xy[1], grid_xy[0], anchor_idx[0][0]])
                updates = updates.write(
                    idx, [box[0], box[1], box[2], box[3], 1, y_true[i][j][4]])
                idx += 1
    return tf.tensor_scatter_nd_update(
        y_true_out, indexes.stack(), updates.stack())

```

In []:

```

def transform_targets(y_train, anchors, anchor_masks, size):
    y_outs = []
    grid_size = size // 32 # suppose we input 416 size then grid size is 416//32 = 13

    # calculate anchor index for true boxes
    anchors = tf.cast(anchors, tf.float32) # casting every anchors to float
    anchor_area = anchors[..., 0] * anchors[..., 1] # calculating the area of anchors
    box_wh = y_train[..., 2:4] - y_train[..., 0:2] # here we are performing xmax-xmin,ymax-ymin using vectors
    box_wh = tf.tile(tf.expand_dims(box_wh, -2),
                      (1, 1, tf.shape(anchors)[0], 1))
    box_area = box_wh[..., 0] * box_wh[..., 1] # these are our Ground Truth Box Area
    intersection = tf.minimum(box_wh[..., 0], anchors[..., 0]) * \

```

```

intersection = tf.minimum(box_wh[..., 0], anchors[..., 0]) * \
    tf.minimum(box_wh[..., 1], anchors[..., 1]) # here we tring to get IOU area
iou = intersection / (box_area + anchor_area - intersection) # simple operation of
Intersection/Union
anchor_idx = tf.cast(tf.argmax(iou, axis=-1), tf.float32) # storing those anchor index which ha
s highest IOU number
anchor_idx = tf.expand_dims(anchor_idx, axis=-1)

y_train = tf.concat([y_train, anchor_idx], axis=-1)

for anchor_idxs in anchor_masks:
    y_outs.append(transform_targets_for_output(
        y_train, grid_size, anchor_idxs))
    grid_size *= 2 # here we are calling the above function for 13*13 grid then, 26*26 grid
then, 52*52 grid.

return tuple(y_outs)

```

Loading YOLOv3 weights into Model

In []:

```

%%time
yolo = YoloV3(classes=80)
yolo.summary()
load_darknet_weights(yolo, "/content/yolov3.weights", False)
yolo.save_weights("/content/yolov3_checkpoint/yolov3.tf")

```

Model: "yolov3"

Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	[(None, None, None, 0		
yolo_darknet (Functional)	((None, None, None, 40620640		input[0][0]
yolo_conv_0 (Functional)	(None, None, None, 5 11024384		yolo_darknet[0][2]
yolo_conv_1 (Functional)	(None, None, None, 2 2957312		yolo_conv_0[0][0] yolo_darknet[0][1]
yolo_conv_2 (Functional)	(None, None, None, 1 741376		yolo_conv_1[0][0] yolo_darknet[0][0]
yolo_output_0 (Functional)	(None, None, None, 3 4984063		yolo_conv_0[0][0]
yolo_output_1 (Functional)	(None, None, None, 3 1312511		yolo_conv_1[0][0]
yolo_output_2 (Functional)	(None, None, None, 3 361471		yolo_conv_2[0][0]
yolo_boxes_0 (Lambda)	((None, None, None, 0		yolo_output_0[0][0]
yolo_boxes_1 (Lambda)	((None, None, None, 0		yolo_output_1[0][0]
yolo_boxes_2 (Lambda)	((None, None, None, 0		yolo_output_2[0][0]
yolo_nms (Lambda)	((None, 116, 4), (No 0		yolo_boxes_0[0][0] yolo_boxes_0[0][1] yolo_boxes_0[0][2] yolo_boxes_1[0][0] yolo_boxes_1[0][1] yolo_boxes_1[0][2] yolo_boxes_2[0][0] yolo_boxes_2[0][1] yolo_boxes_2[0][2]

```

=====
Total params: 62,001,757
Trainable params: 61,949,149
Non-trainable params: 52,608

```

```

CPU times: user 4.02 s, sys: 745 ms, total: 4.76 s
Wall time: 9.46 s

```

Data Parsing Function

In []:

```
def parse_dataset(data, class_dict, size, image, path, yolo_max_boxes, count=0):
    X = []
    Y = []
    for img in tqdm(image):
        x_train = Image.open(path+img) # reading image
        width, height = x_train.size # storing actual width and height so that we can later scale it
        x_train = x_train.resize((size, size)) # resizing
        x_train = np.array(x_train)
        temp_data = []
        # iterating over dataset having info about objects in an image
        for _, row in data[data['img_path']==path+img].iterrows():
            xmin = row.min_w/width
            xmax = row.max_w/width
            ymin = row.min_h/height
            ymax = row.max_h/height
            cls = class_dict[row.category]
            temp_data.append([xmin, ymin, xmax, ymax, cls])
        temp_data = temp_data+[[0,0,0,0,0]]*(yolo_max_boxes-len(temp_data)) # it's like padding
        #return(temp)
        Y.append(temp_data)
        X.append(x_train)
    return(np.array(X), np.stack(np.array(Y)))

# transforming each image and normalizing it in range [0,1]
def transform_images(x, size):
    x = tf.image.resize(x, (size, size))
    x = x/255.0
    return(x)
```

Loading Train and Validation dataset

In []:

```
x,y= parse_dataset(
    df, class_dict, size, train_image[:,], '/content/train/', 116)
# df = data
x = x.astype(np.float32)
y = y.astype(np.float32)

train_dataset = tf.data.Dataset.from_tensor_slices((x,y))

train_dataset = train_dataset.shuffle(buffer_size=64) # Randomizing the data
train_dataset = train_dataset.batch(8) # Setting Batch size
train_dataset = train_dataset.map(lambda x, y: (transform_images(x, size),
                                                transform_targets(y, yolo_anchors, yolo_anchor_masks, size)))
train_dataset = train_dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE) # Prefetching
for faster performance
```

100%|██████████| 2500/2500 [01:39<00:00, 25.10it/s]

In []:

```
train_dataset
```

Out[]:

```
<PrefetchDataset shapes: ((None, 320, 320, 3), ((None, 10, 10, 3, 6), (None, 20, 20, 3, 6), (None, 40, 40, 3, 6))), types: (tf.float32, (tf.float32, tf.float32, tf.float32))>
```

In []:

```
# parsing valid dataset
x,y= parse_dataset(df, class_dict, size, test_image[:,], '/content/train/', 116)
x = x.astype(np.float32)
y = y.astype(np.float32)
val_dataset = tf.data.Dataset.from_tensor_slices((x,y))
```

```

val_dataset = val_dataset.shuffle(buffer_size=16)
val_dataset = val_dataset.batch(8)
val_dataset = val_dataset.map(
    lambda x, y: (transform_images(x, size),
        transform_targets(y, yolo_anchors, yolo_anchor_masks, size)))

val_dataset = val_dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)

```

```
100%|██████████| 873/873 [00:35<00:00, 24.49it/s]
```

In []:

```
val_dataset
```

Out[]:

```
<PrefetchDataset shapes: ((None, 320, 320, 3), ((None, 10, 10, 3, 6), (None, 20, 20, 3, 6), (None, 40, 40, 3, 6))), types: (tf.float32, (tf.float32, tf.float32, tf.float32, tf.float32))>
```

Model Training

In []:

```

model = YoloV3(size, training=True, classes=num_classes)
anchors = yolo_anchors
anchor_masks = yolo_anchor_masks

```

Loading Pretrained model and freezing for transfer learning

In []:

```

%%time
model_pretrained = YoloV3(size, training=True, classes=80)
model_pretrained.load_weights("/content/yolov3_checkpoint/yolov3.tf")
model.get_layer('yolo_darknet').set_weights(
model_pretrained.get_layer('yolo_darknet').get_weights())
freeze_all(model.get_layer('yolo_darknet'))

```

```

CPU times: user 4.39 s, sys: 184 ms, total: 4.58 s
Wall time: 4.5 s

```

Defining Optimizer and Loss for the YOLO model

In []:

```

# we are using graph mode of tensorflow so that we can use our own Gradient Tape
optimizer = tf.keras.optimizers.Adam(lr=1e-4) # Adam optimizers
loss = [YoloLoss(anchors[mask], classes=num_classes) # customized yolo loss define above in utils.
        for mask in anchor_masks]

```

Calculating Loss using Gradient Tape

In []:

```

import datetime
current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
train_log_dir = '/content/320_logs/gradient_tape/' + current_time + '/train' # train dir path
test_log_dir = '/content/320_logs/gradient_tape/' + current_time + '/test' # test dir path
train_summary_writer = tf.summary.create_file_writer(train_log_dir) # train writer
test_summary_writer = tf.summary.create_file_writer(test_log_dir) # test writer

```

In []:

```
avg_loss = tf.keras.metrics.Mean('loss', dtype=tf.float32)
```

```

avg_val_loss = tf.keras.metrics.Mean('val_loss', dtype=tf.float32)
ckpt = tf.train.Checkpoint(step=tf.Variable(1), optimizer=optimizer, model = model)
manager = tf.train.CheckpointManager(ckpt, '/content/320_checkpoints/yolov3_train/tf_ckpts',
max_to_keep=3)
ckpt.restore(manager.latest_checkpoint)
if manager.latest_checkpoint:
    print("Restored from {}".format(manager.latest_checkpoint))
    start = ckpt.step.numpy()
else:
    print("Initializing from scratch.")
    start = 0

epochs = 100
for epoch in range(start, epochs+1):
    for batch, (images, labels) in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            outputs = model(images, training=True)
            regularization_loss = tf.reduce_sum(model.losses)
            pred_loss = []
            for output, label, loss_fn in zip(outputs, labels, loss):
                pred_loss.append(loss_fn(label, output))
            total_loss = tf.reduce_sum(pred_loss) + regularization_loss
            # calculating grads over trainable parameters
            grads = tape.gradient(total_loss, model.trainable_variables) # calculating loss after each
batch
            optimizer.apply_gradients(
                zip(grads, model.trainable_variables)) # then applying gradient optimization on the lo
ss to fine tune the weights

            # writing summary to train file(for tensorboard)

            with train_summary_writer.as_default():
                tf.summary.scalar('avg_loss', total_loss.numpy(), step=epoch)
            # to update avg loss after each batch.
            avg_loss.update_state(total_loss)

        # testing datasets
        for batch, (images, labels) in enumerate(val_dataset):
            outputs = model(images)
            regularization_loss = tf.reduce_sum(model.losses)
            pred_loss = []
            for output, label, loss_fn in zip(outputs, labels, loss):
                pred_loss.append(loss_fn(label, output))
            total_loss = tf.reduce_sum(pred_loss) + regularization_loss

            # writing summary to test file(for tensorboard)

            with test_summary_writer.as_default():
                tf.summary.scalar('avg_val_loss', total_loss.numpy(), step=epoch)
            avg_val_loss.update_state(total_loss)

        # print result
        print("{} , train: {}, val: {}".format(epoch,
            avg_loss.result().numpy(),
            avg_val_loss.result().numpy()))

        ckpt.step.assign_add(1)
        if int(ckpt.step) % 5 == 0:
            save_path = manager.save()
            print("Saved checkpoint for step {}: {}".format(int(ckpt.step), save_path))

        avg_loss.reset_states()
        avg_val_loss.reset_states()

```

Loading Tensorboard

In []:

```

%load_ext tensorboard
%tensorboard --logdir 320_logs

```

Saving the model weights

In []:

```
model.save_weights('/content/Model_Save')
```

Utility functions that help in visualizations

In []:

```
# below function will help in comparing the results when we visualize it after we have pre-trained our model.
```

```
def draw_outputs(img, outputs, class_names):
    boxes, objectness, classes, nums = outputs # predicted outputs
    boxes, objectness, classes, nums = boxes[0], objectness[0], classes[0], nums[0]
    wh = np.flip(img.shape[0:2])
    # iterate through each valid predictions
    for i in range(nums):
        x1y1 = tuple((np.array(boxes[i][0:2]) * wh).astype(np.int32)) #
        x2y2 = tuple((np.array(boxes[i][2:4]) * wh).astype(np.int32))
        img = cv2.rectangle(img, x1y1, x2y2, (255, 0, 0), 2) # it will create a rectangle box
    around object.
    return img
```

```
# below function draws ground truth boxes
class_names = {j:i for i,j in class_dict.items()}
def draw_gt_outputs(path, data, class_names):
    img = plt.imread("/content/train/"+path)
    wh = np.flip(img.shape[0:2])
    nums, classes = [], []
    for _, row in data[data.img_path=="content/train/"+path].iterrows():
        xmin = row.min_w # x1
        xmax = row.max_w # x2
        ymin = row.min_h # y1
        ymax = row.max_h # y2
        nums.append([xmin, ymin, xmax, ymax])
        classes.append(row.label)
    nums = np.array(nums)
    for i in range(nums.shape[0]):
        x1y1 = tuple((np.array(nums[i][0:2])).astype(np.int32)) #
        x2y2 = tuple((np.array(nums[i][2:4])).astype(np.int32))
        img = cv2.rectangle(img, x1y1, x2y2, (255, 0, 0), 2) # it will create a rectangle box
    around object.
    return img
```

Creating new DF with bbox counts

In []:

```
img_data = df['img_path'].value_counts()
df1 = pd.DataFrame({'image':img_data.index, 'count':img_data.values})
```

In []:

```
yolo = YoloV3(classes=num_classes)
yolo.load_weights('/content/Model_Save')
```

Out[]:

```
<tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f65934f4f60>
```

In []:

```
print('HELLO')
```

```
HELLO
```

Final Pipeline

In []:

```
def Final_pipeline(data,indx,visualize=False):
    yolo = YoloV3(classes=num_classes)
    yolo.load_weights('/content/Model_Save')
    class_names = list(class_dict.keys())
    img_raw = tf.image.decode_image(
        open('/content/train/'+data[indx], 'rb').read(), channels=3)
    img = tf.expand_dims(img_raw, 0)
    img = transform_images(img, size)
    img1 = img[0][:]
    t1 = time.time()
    boxes, scores, classes, nums = yolo(img)
    t2 = time.time()

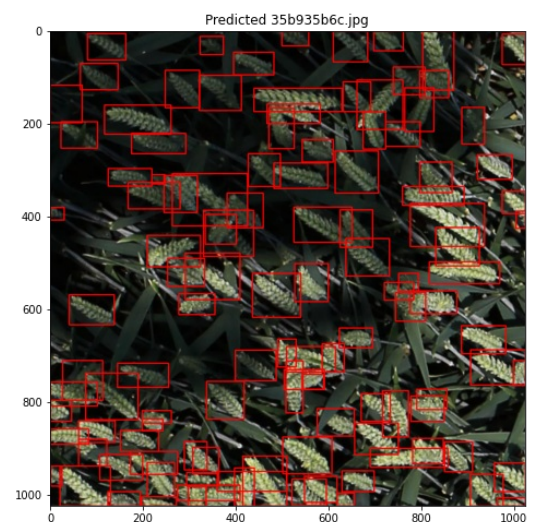

    img = cv2.cvtColor(img_raw.numpy(), cv2.COLOR_RGB2BGR)
    img = draw_outputs(img, (boxes, scores, classes, nums), class_names)


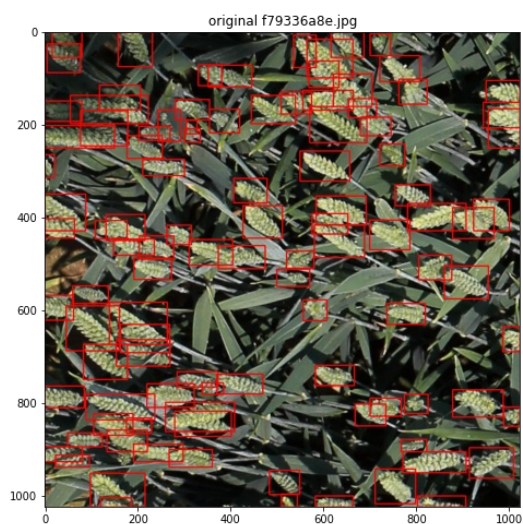

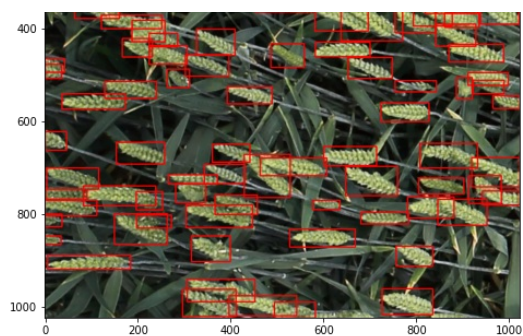
    if visualize:
        plt.figure(figsize=(24,8))
        plt.subplot(1,2,1)
        plt.title("original {}".format(data[indx]))
        gt_img = draw_gt_outputs(data[indx],df,class_names)
        plt.imshow(gt_img)
        plt.subplot(1,2,2)
        plt.title("Predicted {}".format(data[indx]))
        plt.imshow(img)
```

Visualising Dense images

In []:

```
df2 = df1[df1['count']>80]
for i in df2.iloc[0:3,:]['image'].values:
    img_name = i.split('/')[1]
    if img_name in train_image:
        index = train_image.index(img_name)
        Final_pipeline(train_image,index,visualize= True)
    else:
        index = test_image.index(img_name)
        Final_pipeline(test_image,index,visualize= True)
```


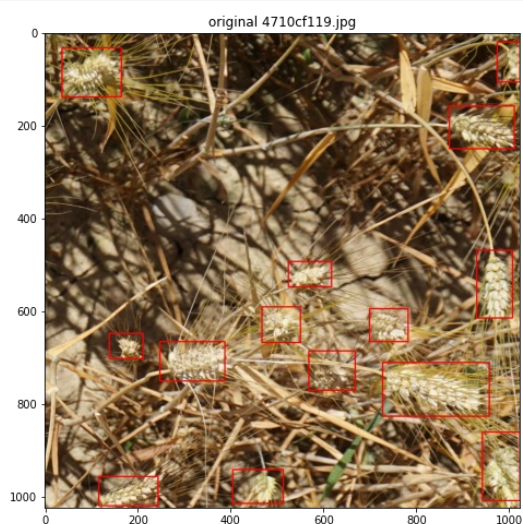


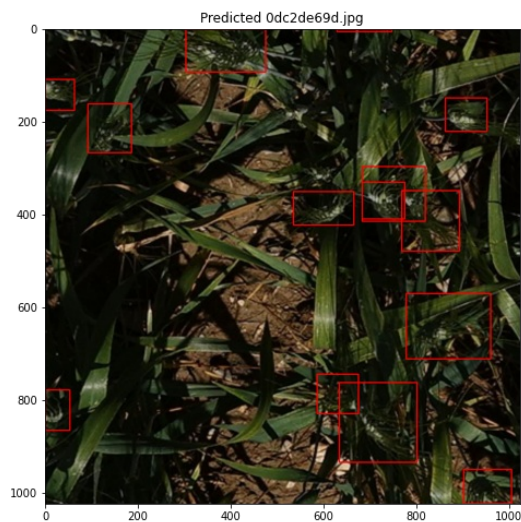
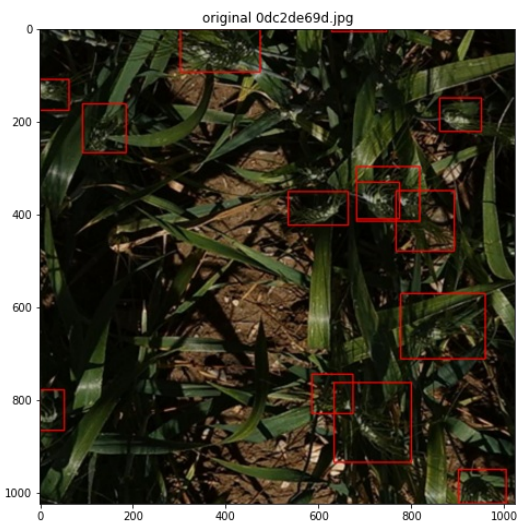




Visualising sparse images

In []:

```
df2 = df1[df1['count']<15]
for i in df2.iloc[0:3,:]['image'].values:
    img_name = i.split('/')[-1]
    if img_name in train_image:
        index = train_image.index(img_name)
        Final_pipeline(train_image,index,visualize= True)
    else:
        index = test_image.index(img_name)
        Final_pipeline(test_image,index,visualize= True)
```


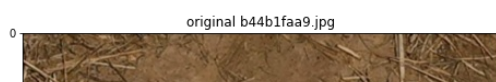

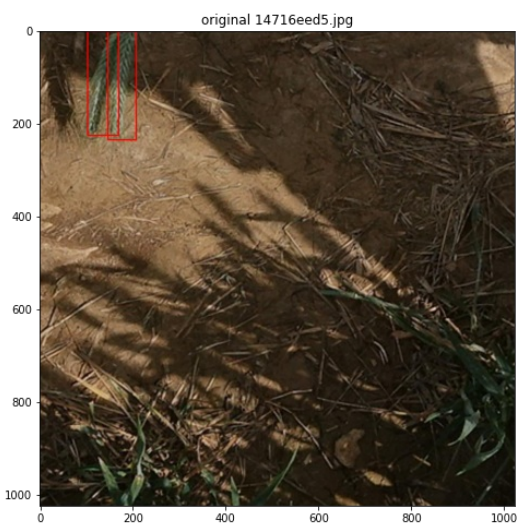






Visualizing very sparse images

In []:

```
df2 = df1[df1['count']<3]
for i in df2.iloc[0:3,:]['image'].values:
    img_name = i.split('/')[-1]
    if img_name in train_image:
        index = train_image.index(img_name)
        Final_pipeline(train_image,index,visualize= True)
    else:
        index = test_image.index(img_name)
        Final_pipeline(test_image,index,visualize= True)
```





Step by Step Procedure

1. Download the gloabl_wheat_detection zip from kaggle and unzip the data.
2. Download the pretrained yolov3 weights.
3. Preprocess the given dataframe by expanding the bbox dimensions and calculate x_min,y_min,x_max,y_max and save it.
4. Split the data into train and validation data.
5. Import all the necessary modules that are needed.
6. Initialize the necessary variables like yolo_anchors, yolo_masks, yolo_max_boxes, batch_size, image_Size, num_Classes, class_dict that are useful for training.
7. Define yolov3 architecture that is implemented completyetly in kears and tensorflow.
8. Define your own custom loss along with the yolo loss for object detection. I used focal loss for my task.
9. Define parse_Dataset function which will help us the map the data that is needed for training by converting it into the format that yolov3 expects.
10. Load both train and validation datasets.
11. Initialize the model and load pretrained weights and freeze the layers that are not required for fine tuning.
12. Start the trainnig. I used gradient tape method to train the model as it gives control over nitty gritty details in the model.
13. While training make sure that we are saving logs and model weights for future use.
14. Once the model is trained save the model weights instead of checkpoints.
15. For making predictions, define yolo architecute again and load the saved weights and make the predictions and visualizations.

Observations

1. When using categorical **cross entropy loss** for negative samples i.e Bounding box for which there is no object is to high and loss for positive samples is too low for which our optimizers tries to lower negative samples loss as it's higher. Therefore, our model wasn't working well at detections.
2. So to overcome the above hurdle we have tried **focal loss** which try to down weight **negative sample loss** and hence improve our model performace.
3. For this model I have taken alpha parameter of **focal loss 0.85**, I have also tried 0.60,0.65,0.70,0.75,0.80,0.85 but at last

0.75 suited best for this model.

Future Scope or Work

1. In our model we have used YOLOv3 which is a really good object detection technique but at the time of making this case study we already have YOLOv5 which is state of the art model.
2. So using YOLOv5 might get some better results as yolov5 is faster than yolov3 and its accuracy is reasonable and comparable too.
3. We have only 3k images for training, and we all know that more the data more the model performance for deep learning models. So training to incorporate Augmentations might be helpful.
4. There is a module called Albumentations that will help immensely while augmenting images with bounding boxes.
5. Instead of yolo we can use other models like Faster RCNN which is build for object detection tasks.
6. We can get some more images from internet and pseudo labelling them carefully and adding those images for training might help too.

What I Gained?

1. By doing this case study I got to know every single thing that YOLOv3 does internally for object detections.
2. By implementing YOLOv3 completely in keras or tensorflow helped me gain confidence in me so that anyone with good knowledge can build an research on their own using all the tools available.
3. If loss function is continuous and differentiable we can always back propagate the loss and minimize loss using the techniques we learned in deep learning.
4. Finally I think i will be able to implement my own versions of other models in keras or tensorflow.